

Best Practices for Migrating MySQL to the Cloud

Juan Pablo Arruti
Percona



Agenda

- IaaS vs DBaaS
- Migrating Data
- Replication between On-Premises and Cloud
- Testing Cloud Environments
- High Availability
- Monitoring
- Backups



laaS vs DBaaS

What is IaaS?

Infrastructure as a Service (IaaS)

- Fundamentals Compute Resources
 - Servers, Storage, Network
- Provisioned and managed over the internet
- Complete control of servers
- AWS Services
 - EC2, EBS, VPC

And DBaaS?

Database as a Service (DBaaS)

- Provides database service (Instance or Cluster)
- Targeted to easily
 - Setup, Operate, Scale
- Manages common administration tasks
 - Backups, Patching, Failure Detection, Failover
- No OS access
- No Super privilege
- AWS
 - **RDS**, DynamoDB, Redshift

Pros and Cons of IaaS

Pros

- More control and flexibility
- Wide Instance Types
- Cheaper than RDS

Cons

- More operational work



Pros and Cons of DBaaS

Pros

- Easy to Manage
- Less operational tasks

Cons

- Less control and flexibility
- More expensive
- Limited Instance types

Which Do You Choose?

IaaS

- Database needs specific tuning or feature
- Available resources for operational tasks

DBaaS

- Need focus on data and code
- Generic setups are okay

Migrating Data

Best Practices

- Make it simple
- Migrate to same or higher minor version
- Avoid major version upgrades

Migrating Data to IaaS

- Similar to on-premises databases
- Use Physical Backups for large databases
 - XtraBackup, Cold Backups
- Logical Backups for small databases



Migrating Data to DBaaS

- Logical Backups
 - Access through MySQL Client
 - mysqldump, mysqlpump, mydumper
- Physical Backups are possible
 - XtraBackup*

Only available for AWS RDS

Migrating Data to DBaaS

Available MySQL Client Tools

- **mysqldump**
 - Most adopted tool
 - Single-threaded
- **mysqlpump**
 - Introduced in MySQL 5.7
 - Parallel backups
 - Restores are Single-threaded
- **mydumper/myloader**
 - Parallel backups and restores



Migrating Data to DBaaS

Best Practices for MySQL Clients

- Export all objects first
 - `--no-data --routines --events --triggers`
- Then export only data
 - `--no-create-info --no-create-db`
 - `--routines=no --events=no --triggers=no`
- Enable `log_bin_trust_function_creators` if `log_bin=1`
- Change object definer
 - `DEFINER='user'@'host'`
- Force load and check all errors
 - `--force`



Migrating Data to RDS

Best Practices for MySQL Clients

- Increase `max_allowed_packet` (Default 4 MB)
- `time_zone` can be modified in parameter group (Default UTC)
 - RDS uses mysql schema Time Zone Tables
 - Recommended
 - Set session `time_zone` to match source database

Migrating Data to RDS

Speeding Up Logical Restore

- EC2 and RDS in same AZ
- Disable Multi-AZ
- Increase IOPS
- Modify Default Settings
 - Relax Durability
 - `sync_binlog != 1`
 - `innodb_flush_log_at_trx_commit != 1`
 - Tune InnoDB
 - Increase `innodb_log_file_size` (Default 128 MB)
 - Increase `innodb_buffer_pool_size` (Default $\text{DBInstanceClassMemory} * 3/4$)

Restore Amazon RDS from Xtrabackup

- Overview
 - Take backup from database
 - Upload into S3 bucket
 - Create new instance from the backup
 - Amazon MySQL RDS
 - Amazon Aurora MySQL

Restore Amazon RDS from Xtrabackup

- Limitations
 - Supported MySQL 5.6 and 5.7
 - Source/Target major versions must match
 - Target minor version must be higher
 - Source tables defined within default datadir
 - 6 TB database size limit
 - Source database can't be encrypted
 - User accounts, functions, stored procedures and time zone info are not imported automatically

Replication Between On-Premises and Cloud

Replicating to the Cloud

- IaaS
 - Same as replication in on-premises
- DBaaS
 - Implementation and its limitations depends on the cloud provider



Replicating to the Cloud

Best Practices

- If latency is high
 - Use compression for Master/Slave protocol
 - `slave_compressed_protocol=1`
 - Monitor replication lag with **pt-heartbeat**
- Ensure tables have Primary Key
 - `binlog_format = ROW`

External Master on AWS RDS

- Easy to set
- NO `binlog_format` constraints (MIXED, ROW, STATEMENT)
 - Recommended ROW to avoid `time_zone` mismatch
- Log File Position or GTID based
- No filtered Replication is allowed
- Replication administration using procedures
 - `mysql.rds_set_external_master`
 - `mysql.rds_start_replication`
 - ...

Testing Cloud Environments

Why Benchmark Cloud?

- Cloud resources may not map directly
- Validate if cloud instance is able to handle traffic
- Choose between IaaS and DBaaS
- Available tools
 - sysbench, pt-upgrade, **Query Playback**, **ProxySQL**

Query Playback

Key aspects

- Percona Labs GitHub repository
 - No active development
- Executes Queries in logs
 - Slow Query log, General log
- Compares execution results with Log
- Servers data should be consistent
- NO read-only option
- Multi-threaded
 - Queries executed at arrival time

Query Playback

Example Report

```
SELECT c FROM sbtest37 WHERE id=505; -->  
thread 67 slower query was run in 86 microseconds instead of 34
```

Detailed Report

SELECTs : 1858522 queries (19297 faster, 1839225 slower)
...

Report

Executed 2161872 queries
Spent 00:09:08.631886 executing queries versus an expected 00:04:43.697328 time.
23610 queries were quicker than expected, 2138262 were slower
A total of 0 queries had errors.
Expected 40606531 rows, got 40606533 (a difference of 2)
Number of queries where number of rows differed: 2.

Average of 113782.74 queries per connection (19 connections).



ProxySQL

What is ProxySQL?

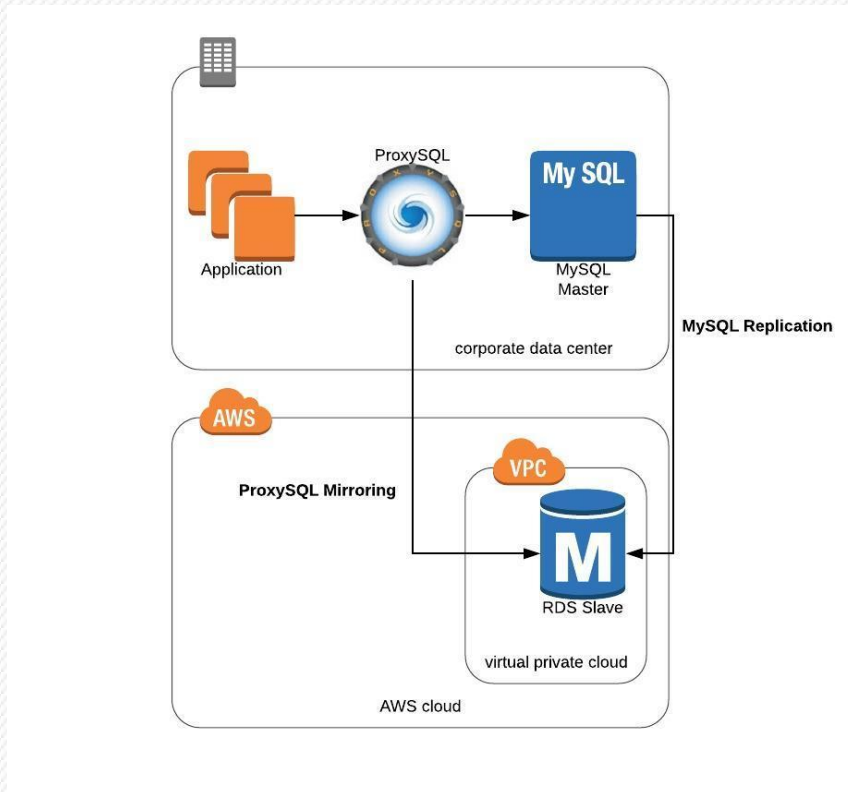
- GPL High Performance MySQL Proxy
- MySQL Protocol Aware
 - Clients connect to ProxySQL
 - Requests are evaluated
 - Actions are performed
 - Routing, Re-write, **Mirroring**

ProxySQL Query Mirroring

How does it work?

- Each client executes a query
- ProxySQL receives each query
- Query Processor identifies if the query is Mirrored
- Associates the Query to a Thread Pool
- Executes each Query in the Pool

ProxySQL Query Mirroring



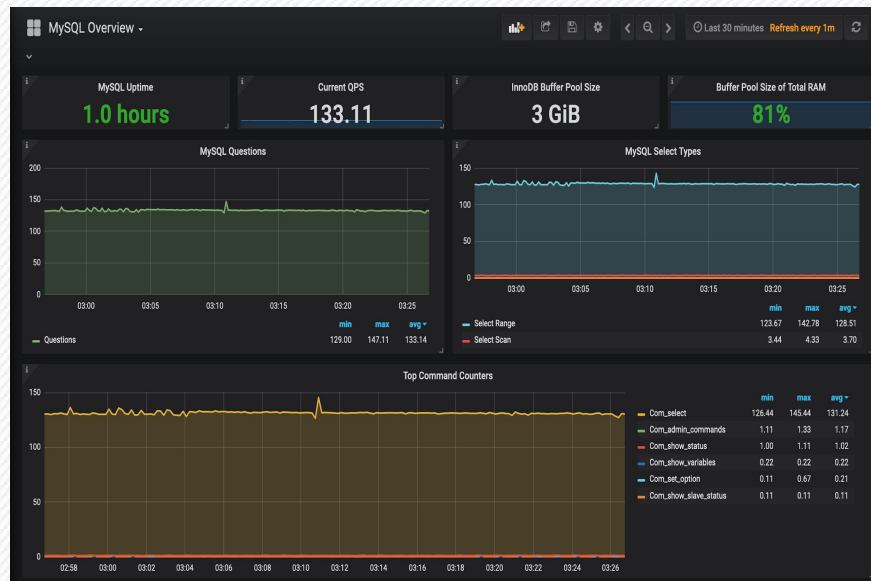
ProxySQL Query Mirroring

Limitations

- May execute Queries in different order
- Some Queries may not be executed
- Data consistency is not guaranteed
- Adds load to ProxySQL process
- Prepared statements are not supported
- No report is provided



Query Mirroring Example



Query Mirroring Example

```
mysql> select * from stats_mysql_query_digest where digest_text like 'SELECT c FROM sbtest1 %' ORDER BY  
hostgroup \G
```

```
***** 1. row *****
```

```
  hostgroup: 1
```

```
  schemaname: sbtest
```

```
  username: jptest
```

```
    digest: 0x290B92FD743826DA
```

```
digest_text: SELECT c FROM sbtest1 WHERE id BETWEEN ? AND ?
```

```
count_star: 48864
```

```
first_seen: 1558761934
```

```
last_seen: 1558765725
```

```
sum_time: 14788745778
```

```
min_time: 2877
```

```
max_time: 3733095
```

```
***** 2. row *****
```

```
  hostgroup: 2
```

```
  schemaname: sbtest
```

```
  username: jptest
```

```
    digest: 0x290B92FD743826DA
```

```
digest_text: SELECT c FROM sbtest1 WHERE id BETWEEN ? AND ?
```

```
count_star: 48832
```

```
first_seen: 1558761936
```

```
last_seen: 1558765725
```

```
sum_time: 16477562501
```

```
min_time: 2786
```

```
max_time: 4651554
```



High Availability

Load Balancers

Why are they useful?

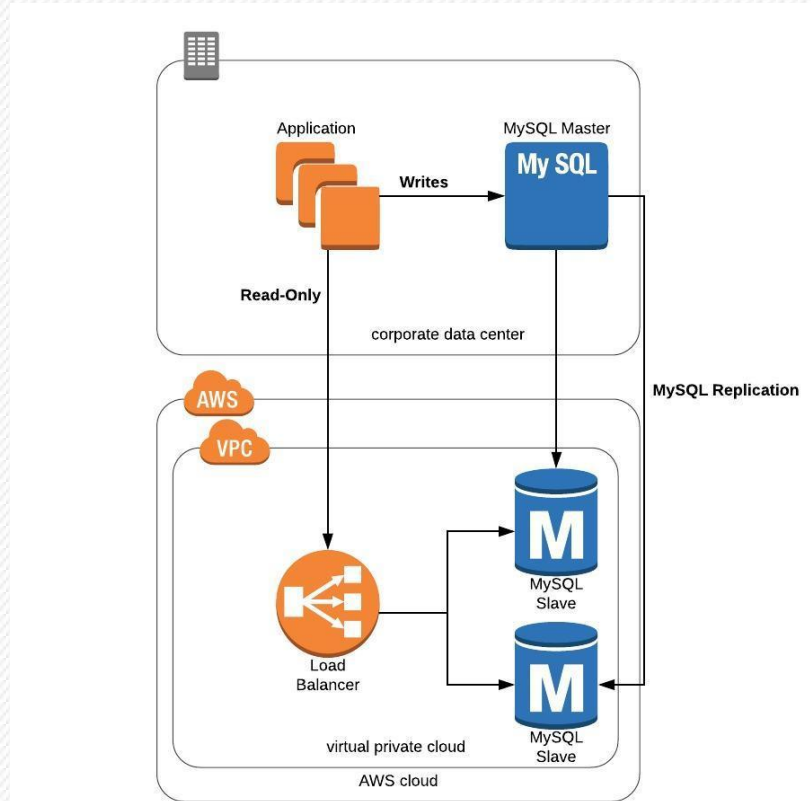
- Phased migration approach
- Routing read-only traffic
- Monitor the new environment
- Adjust it if necessary
- Confirm the environment is stable
- Minimize downtime
- Avoid modifying App connection string
- Available Load Balancers
 - **Cloud Load Balancing**
 - **ProxySQL**



Cloud Load Balancing

- TCP Load Balancer
- Distributes traffic
- AWS Elastic Load Balancer
 - RDS not supported
 - Routes traffic only across Availability Zones
 - Custom MySQL health checks can be used

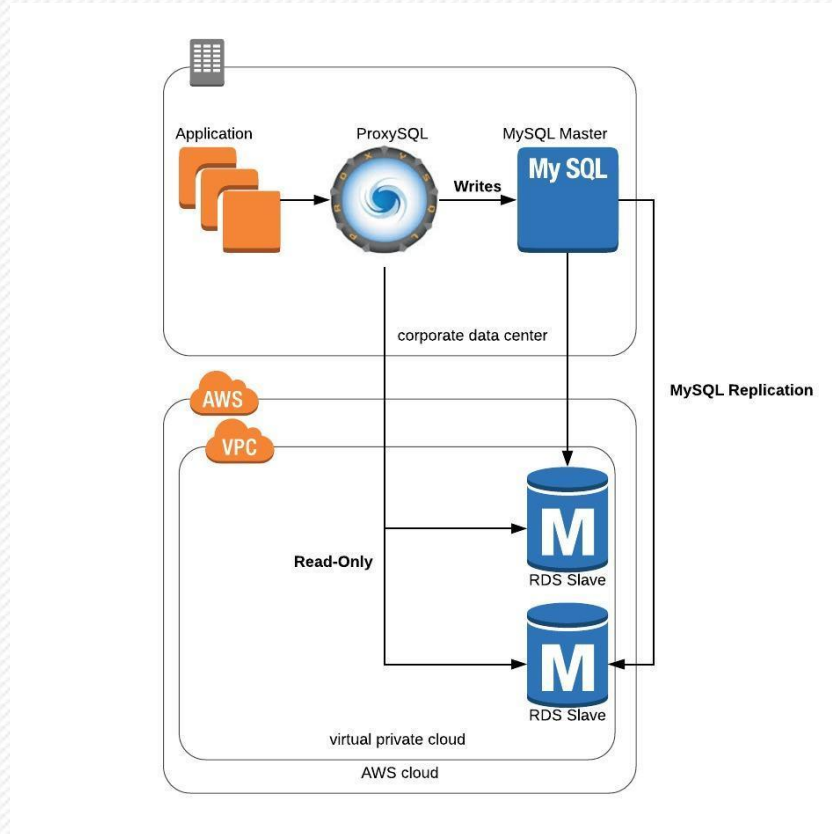
AWS ELB - Phased Migration



ProxySQL

- SQL Connection Endpoint
- Routes traffic to MySQL databases
- Splits read/write traffic
 - Based on `read_only` value
- Supports DBaaS

ProxySQL - Query Routing



DBaaS High Availability

How is it achieved?

- Data redundancy
- Automatic failover

How is it implemented?

- Each cloud vendor implements it differently

AWS RDS - Multi-AZ

Key aspects

- Synchronous Standby Replica - DRBD
- Block replicated to Different Availability Zone than Primary
- Secondary is used for backups
- Failover takes place by internal DNS change

Limitations

- Reads on Secondary are not possible
- DML Overhead



AWS RDS Read-Replicas

Key aspects

- MySQL asynchronous replication
- Scale-out Reads
- Promote to stand-alone database
- Within Same AZ, Cross AZ, Cross Region
- Easy to implement



Monitoring

Monitoring MySQL in the Cloud

Best practices

- Establish a baseline
- Measure workload under different conditions
- Compare cloud instances and on-premises

Available tools

- Cloud Monitoring System
- Percona Monitoring and Management

AWS CloudWatch

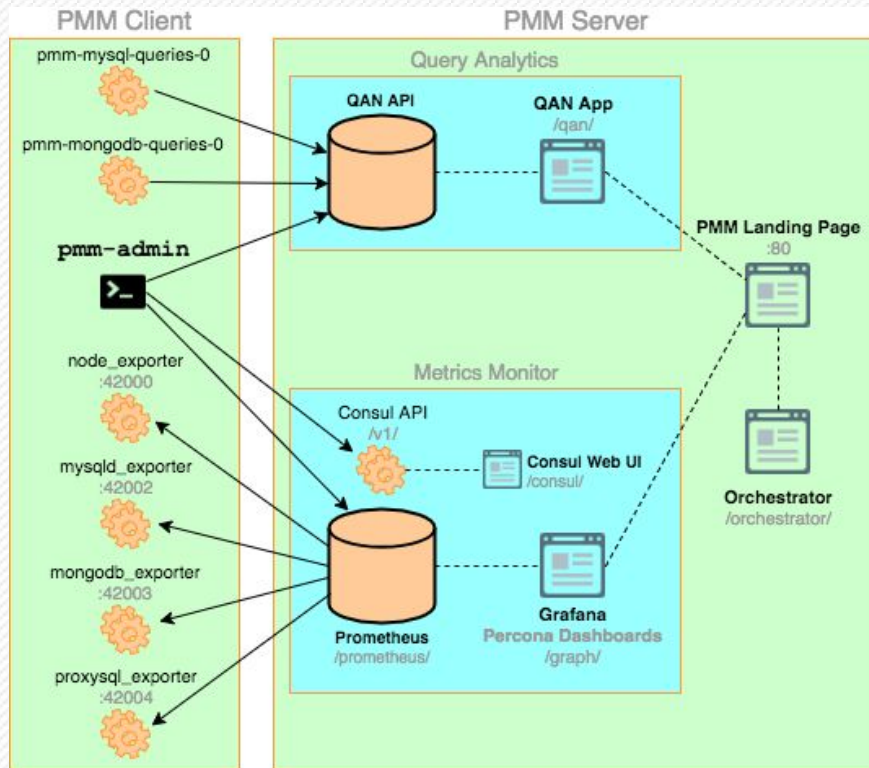
What is CloudWatch?

- Monitors AWS resources
- Metrics collected automatically
- Custom dashboards
- Create alarms and send notifications based on metrics
- Launch additional resources or stop them
- High-level Database metrics
 - BinLogDiskUsage
 - DatabaseConnections
 - ReplicaLag

Percona Monitoring and Management

- Open source monitoring platform
- MySQL, MongoDB, PostgreSQL and ProxySQL
- Detailed metrics for DB Servers
- Query Analytics shows current queries and stats
- Supports AWS MySQL RDS and Aurora

PMM Architecture



Monitoring RDS in PMM

Recommendations

- PMM Server in same AZ than RDS
- Don't use T2 instances
- Use Elastic IP address
- Don't use RDS admin user



Monitoring RDS in PMM

- Enable `performance_schema` for Query Analytics
- Enable enhanced monitoring
 - OS level metrics

Monitoring

Enhanced monitoring

☒ Enable enhanced monitoring
Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU.

☐ Disable enhanced monitoring

Monitoring Role

rds-monitoring-role ▼

Granularity

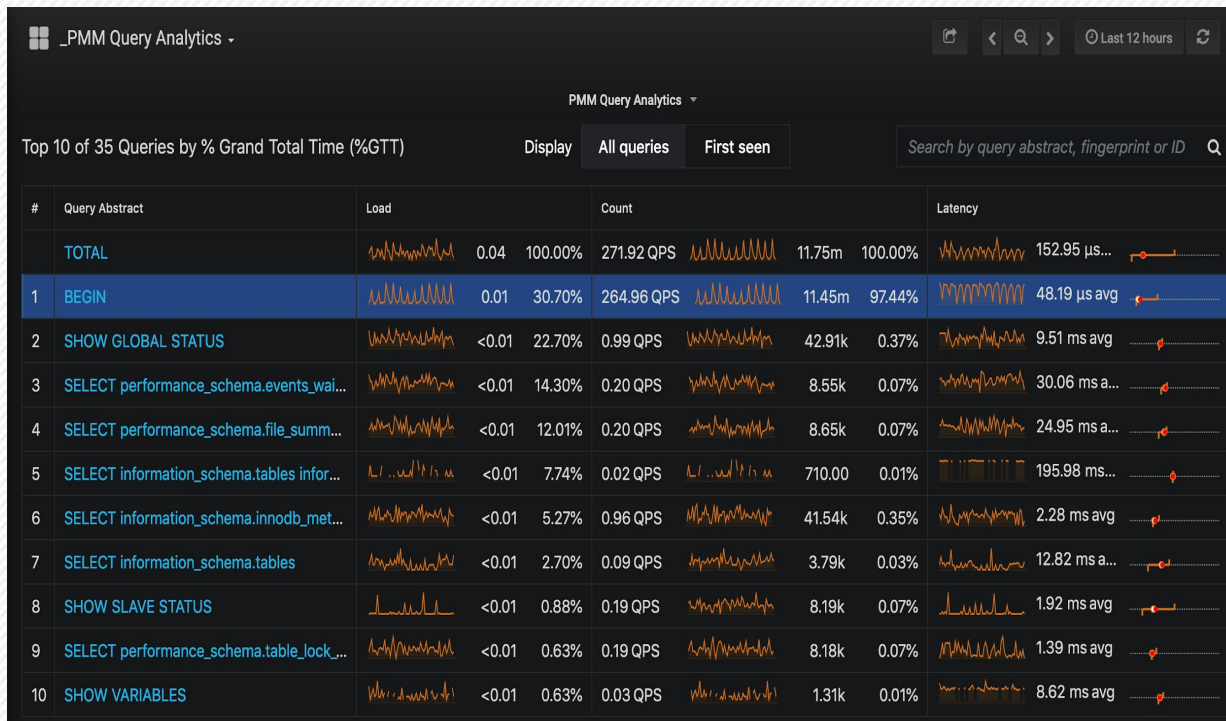
60 seconds ▼



PMM MySQL Overview Dashboard



PMM Query Analytics



Backups

Cloud Storage

Key aspects

- Object storage - AWS S3
- Designed for durability - 99.999999999%
- Highly available
- Scalable
- Secure
- Range of storage classes
 - Standard, Infrequent Access, Glacier
 - Lifecycle policies
- Off-site storage



MySQL Backups to AWS S3

mysqldump - example

- Copy dump file with aws cli

```
shell> mysqldump --all-databases > mysql.dump.sql
```

```
shell> gzip mysql.dump.sql
```

```
shell> aws s3 cp mysql.dump.sql.gz s3://mysqlbucketprod/mysql.dump.sql.gz  
upload: ./mysql.dump.sql.gz to s3://mysqlbucketprod/mysql.dump.sql.gz
```

or

```
shell> mysqldump --all-databases | gzip | aws s3 cp - \  
s3://mysqlbucketprod/mysql.dump.sql.gz
```

MySQL Backups to AWS S3

XtraBackup - example

- Stream files with **xbstream**
- Upload stream with **xbcloud**

```
shell> xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp \
      --target-dir=/tmp | \
      xbcloud put --storage=s3 \
      --s3-endpoint='s3-us-west-2.amazonaws.com' \
      --s3-access-key='AKIAJ6HPUNXNZPTL2AAA' \
      --s3-secret-key='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' \
      --s3-bucket='mysqlbucketprod' \
      --parallel=10 \
      $(date +"%Y%m%d%H%M%S")-full_backup
```

RDS Backups

Automated backups

- Defined backup window
- Stored in S3
- InnoDB tables
- Elevated latency and IO freeze - except for Multi-AZ
- Point in time recovery
- Retention period = 1 - 35 days

Database snapshots

- Manually initiated
- No point in time recovery

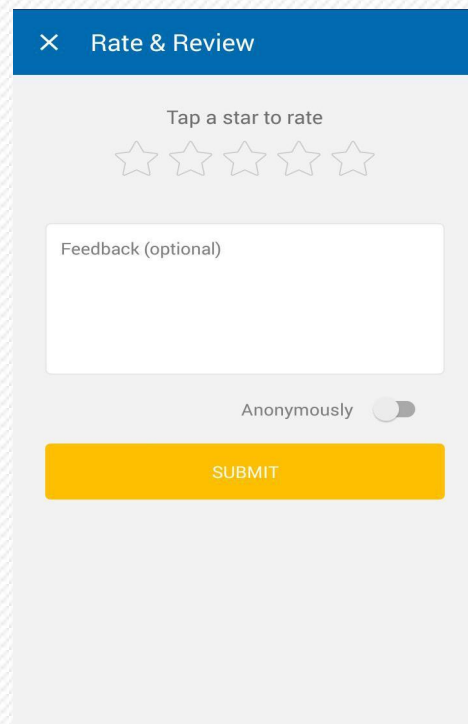
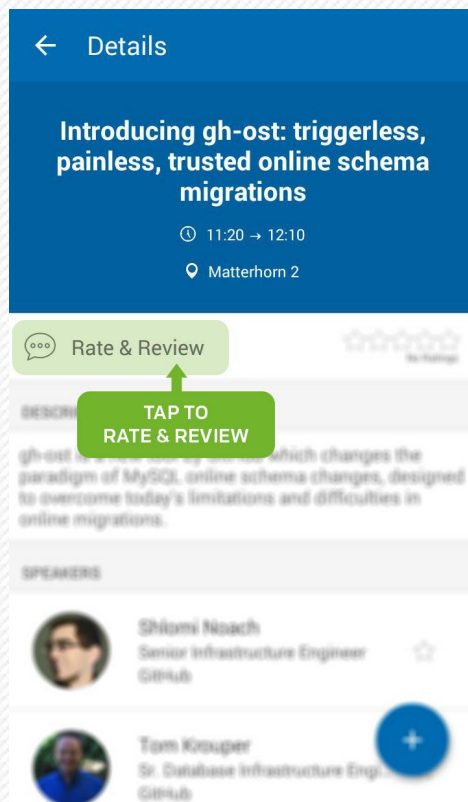
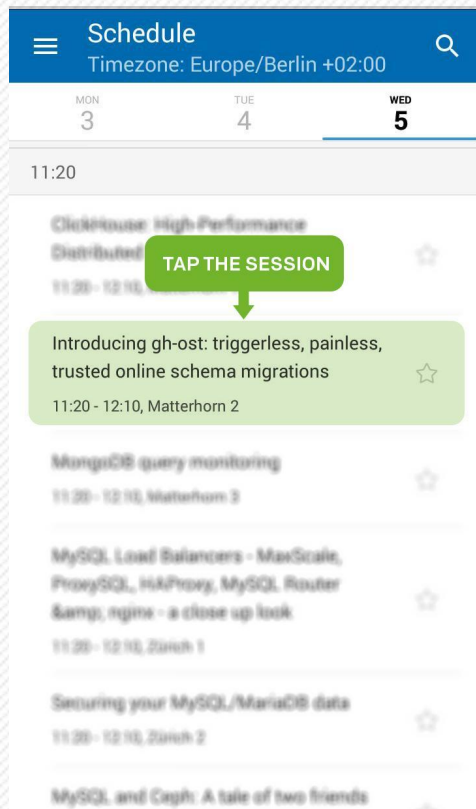
Questions

Thank You to Our Sponsors



PERCONA
LIVE

Rate My Session



Thank you!
