




TRACING POD²POD NETWORK TRAFFIC in KUBERNETES





When you deploy a Pod, the following things happen:


1. The pod gets its own **network namespace**.
2. An **IP address** is assigned.
3. Any containers in the pod **share the same networking namespace** and can see each other on localhost.




For a pod to reach other pods it must first have access to the node's root namespace.

This is achieved using a virtual eth pair connecting the 2 namespaces: pod and root.

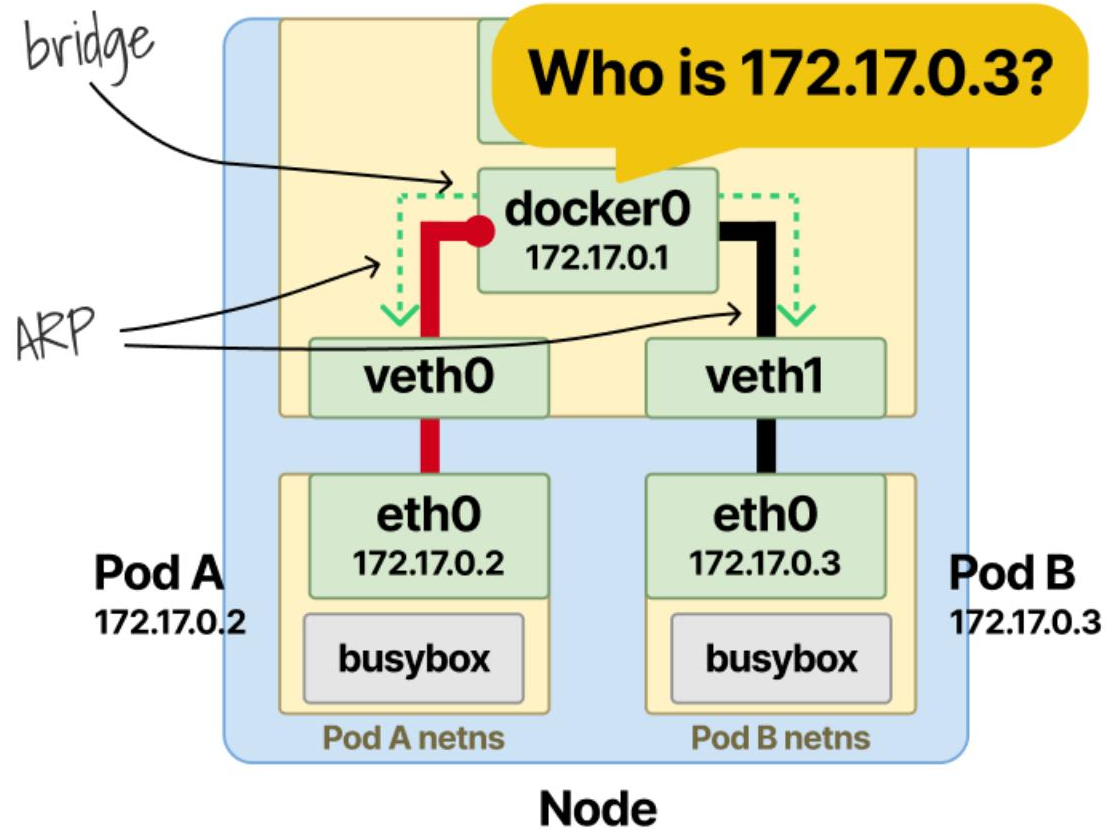
The bridge allows traffic to flow between virtual pairs and traverse through the common root namespace.



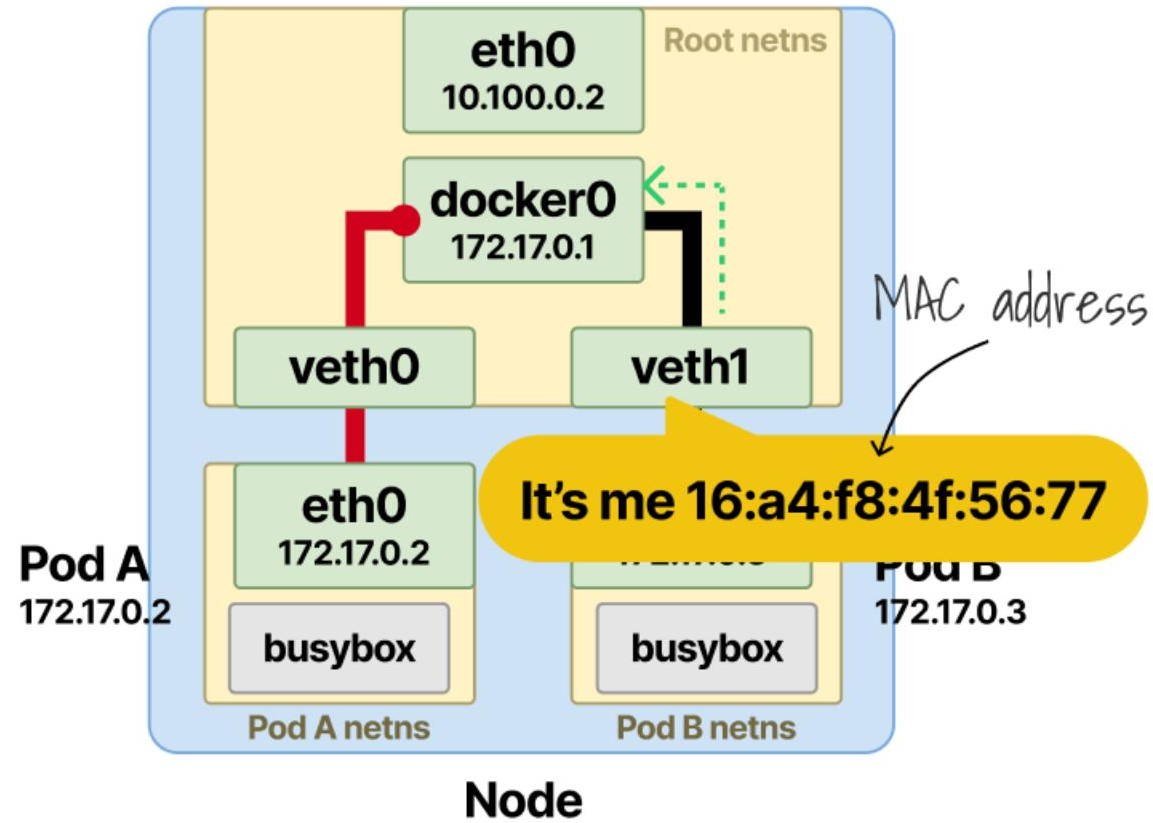
So what happens when Pod-A wants to send a message to Pod-B?
Since the destination isn't one of the containers in the namespace Pod-A sends out a packet to its default interface `eth0`. This interface is tied to the veth pair and **packets are forwarded to the root namespace.**




The ethernet bridge, acting as a virtual switch, has to somehow resolve the destination pod IP (Pod-B) to its MAC address.




The ARP protocol comes to the rescue.
An ARP broadcast is sent on all connected devices when the frame reaches the bridge.
The bridge shouts "Who has Pod-B IP address?"



A reply is received with the MAC address of the interface that connects Pod-B, then this information is stored in the bridge ARP cache (lookup table).



Once the IP and MAC address mapping is stored, the bridge looks up in the table and forwards the packet to the correct endpoint. The packet reaches Pod-B veth in the root namespace, and from there, it quickly reaches the eth0 interface inside the Pod-B namespace.



With this, the communication between Pod-A and Pod-B has been successful.

For pods that need to communicate across different nodes, an additional hop is required.

If you are interested in learning more, check out this article: <https://learnk8s.io/kubernetes-network-packets>



That's all Folks!

That's it! Thanks for reading this far!

I post deep-dives like this on a regular basis. Follow me if you don't want to miss the next story!

